

PLSC 197 / SODA 110N: Social Data, Technology, and Artificial Intelligence

Fall 2026

Professor Christopher Zorn
Department of Political Science
Pennsylvania State University
Office: Susan Welch Liberal Arts Building, Room 545
E-mail: zorn@psu.edu
Mondays, Wednesdays, and Fridays, 1:25-2:15 p.m.
[Health and Human Development Building](#), Room 350

Course Description

“How do everyday traces of human activity become the building blocks of artificial intelligence (AI)? This course invites students to explore the world of AI and digital technology from the ground up, beginning with the data we create, and following its journey into powerful algorithmic systems that shape society. Students will unpack how human-generated data is collected, structured, and transformed into machine learning models, discovering both the technical and human sides of AI. Through real-world case studies, interactive projects, and critical discussions, students will gain hands-on experience with data systems and analytics tools while examining the social, political, and ethical questions that surround AI in practice. Whether you’re curious about technology, policy, design, or social impact, this course offers an accessible and thought-provoking introduction with no programming experience required.”

That’s the course description, from the Penn State website. It was, naturally, generated by AI (and me).

PLSC 197 / SODA 110N satisfies the following:

- PLSC B.A. majors can use this course toward the 100-level or above requirement and the requirement for any political science or related area.
- PLSC (B.A. or B.S.) majors can receive GS credit if they register for the course as SODA 197N, since PLSC majors must take GS credits outside of the discipline of political science.
- SoDA majors can use this course toward the six credits of additional political science requirement.

All course materials (including this syllabus) will be available on a dedicated Github repo, which can be found at: <https://github.com/PrisonRodeo/PLSC197-FA2026-git>. Throughout this syllabus, hot links are in [Penn State Blue](#).

Things To Read, Etc.

There is (currently) no text for this course. All of the required and recommended readings (and other materials) assigned below will be made available on the course [github repository](#). Note as well that (a) this is an introductory-level course, (b) this course is in an area that is changing quickly, and (c) as a result, no

one course (or even reading list) can hope to be comprehensive. There are many, many books (and other resources) out there for deeper treatments of data analysis and visualization, machine learning, artificial intelligence, and/or digital ethics.

Grading and Assessment

For this course, grades will be assessed based on a total of 1000 possible points, to be distributed as follows:

- Course attendance / participation: 100 points
- Two short assignments: 100 points each
- Two applied labs: 150 points each
- Final examination: 400 points

Short assignments will be available for completion weekly; the two applied labs will take place the weeks of September 14-18 and October 26-30, 2026. The date and time of the final exam will be announced later in the course. Final grades will be assigned following the [grading scheme for IST World Campus](#). More details about each assignment will be made available in class, and on the course Github repo.

A Few Other Useful Resources

None of these are requirements, but you may find them useful:

- [Hugging Face](#) is the place where AI (really) happens.
- [Towards Data Science](#) is always a good read.
- [Wired](#) is a bit pop-culture-y, but remains a good, reliable (and cheap) place to keep abreast of news in the AI community.
- A small sample of available AI tools (August 2026):
 - [ChatGPT](#)
 - [Claude](#)
 - [DeepSeek](#)
 - [Google Gemini](#)
 - [Microsoft Copilot](#)
 - [MetaAI](#)
 - [Grok](#)
 - [Adobe Firefly](#)
 - [Cursor](#)
 - [OpenArt / Dall-E 3](#)
 - [MidJourney](#)
 - [Luma](#)
 - [Runway](#)
 - [Kling](#)
 - [Suno](#)
 - [Udio](#)
 - [Elicit](#)

Obligatory Statement on Academic Integrity

Penn State defines academic integrity as “the pursuit of scholarly activity in an open, honest and responsible manner” (Senate Policy 49-20). Dishonesty of any kind will not be tolerated in this course. Dishonesty includes, but is not limited to, cheating, plagiarizing, fabricating information or citations, facilitating acts of academic dishonesty by others, having unauthorized possession of examinations, submitting work of another person or work previously used without permission from the instructor, or tampering with the academic work of other students. Students facing allegations of academic misconduct should not drop the course; those who do will be added to the course again and will be expected to complete course work and meet course deadlines. If the allegations are dismissed, then the drop will be permitted. Students found responsible for academic misconduct often receive academic sanctions, which can be severe, and put themselves at risk for disciplinary sanctions assigned by the University’s Office of Student Conduct (see [Senate Policy G-9](#)).

In cases of any violation of academic integrity it is the policy of the Department of Political Science to follow procedures established by the College of the Liberal Arts. More information on academic integrity and procedures followed for violation can be found [here](#).

Obligatory Statement on Accommodations for Disabilities

Penn State welcomes students with disabilities into the University’s educational programs. Every Penn State campus has an office for students with disabilities. Student Disability Resources (SDR) website provides [contact information for every Penn State campus](#). For further information, please visit the [Student Disability Resources website](#).

In order to receive consideration for reasonable accommodations, you must contact the appropriate disability services office at the campus where you are officially enrolled, participate in an intake interview, and provide documentation: [see documentation guidelines](#). If the documentation supports your request for reasonable accommodations, your campus disability services office will provide you with an accommodation letter. Please share this letter with your instructors and discuss the accommodations request with them as early as possible. You must follow this process for every semester that you request accommodations.

Students with disabilities participating in internship, practicum, student teaching, or other experiential learning opportunities as part of their degree requirements may also be eligible for reasonable accommodations to ensure equal access and opportunity. These accommodations are determined through an interactive process involving the student, their University supervisor, and the site supervisor. Student Disability Resources can assist students with identifying potential barriers, facilitating accommodation requests, and coordinating with University supervisors to promote inclusive learning experiences.

Obligatory Statement on Counseling and Psychological Services

Many students at Penn State face personal challenges or have psychological needs that may interfere with their academic progress, social development, or emotional wellbeing. The university offers a variety of confidential services to help you through difficult times, including individual and group counseling, cri-

sis intervention, consultations, online chats, and mental health screenings. These services are provided by staff who welcome all students and embrace a philosophy respectful of clients' cultural and religious backgrounds, and sensitive to differences in race, ability, gender identity and sexual orientation.

Counseling and Psychological Services at University Park (CAPS)
(<http://studentaffairs.psu.edu/counseling/>): 814-863-0395

Counseling and Psychological Services at Commonwealth Campuses
(<http://senate.psu.edu/faculty/counseling-services-at-commonwealth-campuses/>)

Penn State Crisis Line (24 hours / 7 days/week): 877-229-6400. Crisis Text Line (24 hours / 7 days/week): Text LIONS to 741741.

Obligatory Statement on Educational Equity and Reporting Bias

Penn State takes great pride to foster a diverse and inclusive environment for students, faculty, and staff. Consistent with University Policy AD29, students who believe they have experienced or observed a hate crime, an act of intolerance, discrimination, or harassment that occurs at Penn State are urged to report these incidents as outlined on the University's Report Bias webpage (<http://equity.psu.edu/reportbias/>).

Non-Obligatory Statement on Extended Absences

During your enrollment at Penn State, unforeseen challenges may arise. If you ever need to miss an extended amount of class in such a circumstance, please notify your professor so you can determine the best course of action to make up missed work. If your situation rises to a level of difficulty you cannot manage on your own with faculty support, reach out to the Student Care & Advocacy office by phone at (814)-863-2020 or email them at StudentCare@psu.edu.

Non-Obligatory Statement on Religious Observances

The [Religious and Spiritual Observances Calendar](#) is compiled by the Center for Spiritual and Ethical Development in consultation with campus and community religious leaders. It specifies those holy days of the major world religions for which observance may require students to depart from their normal routine at the University. Please note that only those holy days which occur when Penn State classes are in session are listed. This is not, therefore, an exhaustive list of all major holy days in each religious tradition. If you have any questions regarding this list, or the dates of other holy days not listed, please contact the Center for Spiritual and Ethical Development.

Non-Obligatory Statement on Generative AI, Large Language Models, etc.

It's unlikely you be in this course if you weren't well aware of the existence of large language models (LLMs) – e.g., [ChatGPT](#) – and other artificial intelligence (AI) tools for language / image creation. Having been described as everything from [making everyone their own version of Tony Stark](#) to a [Lovecraftian](#)

shoggoth (and most things in between), LLMs are currently creating a sometimes-depressing, sometimes-hilarious panic among faculty in legacy academic disciplines and programs. As we'll discuss, most of that panic revolves around the use of LLMs to "cheat," in the traditional sense: to create work that deceptively gives the impression that the student knows something they do not. Beyond its intrinsically duplicitous nature, such use in a conventional classroom setting gives rise to concerns about equity and (potentially) devalues the experience / credential for other class members.

The other side of the LLM equation is that they are powerful tools for augmenting learning and creating new knowledge. Experience suggests that, in line with other technological advances (the printing press, personal computers, search engines, etc.), it is wiser to adapt to LLMs than to attempt to limit or ban their use. This is especially true in a course like this one, where restricting or banning AI from class would be even more ironic than it would be counterproductive. Accordingly, enrollees in PLSC 197 / SODA 110N are welcome to use generative AI tools, such as ChatGPT, to assist them with their work in the course. In doing so, it is important to remember that such AI tools are capable of making errors, and that it is each student's responsibility to verify the information they receive from the such a tool. In addition, any information obtained from a generative AI source must be noted/cited in the student's work, just as they would cite any other source.

Course Schedule

Daily reading assignments are listed for each class session. Note that reading assignments are listed according to the day on which the subject matter will be discussed; they should therefore be read prior to that date.

I. Data

August 24: Course Introduction

- No readings assigned.

August 26: Human Behavior as Data

Read:

- Mimi Onuoha. The Point of Collection. *Data & Society*, February 2016. URL <https://datasociety.net/points/the-point-of-collection/>.
- Nicola Jones. The AI revolution is running out of data. What can researchers do? *Nature*, 636(8042): 290–292, December 2024. URL <https://www.nature.com/articles/d41586-024-03990-2>.

August 28: Introduction to Penn State's AI Studio

- No readings assigned.

August 31: Data Structures and Formats

Read:

- Wuraola Ademola-Shanu. What Is Data? Structured vs. Unstructured Data (Explained Simply for Beginners). *Medium*, August 2025. URL <https://medium.com/@TheWuraolah/what-is-data-structured-vs-unstructured-data-explained-simply-for-beginners-88c46fc9c28f>, (and/or)
- Keith D. Foote. Structured vs. Unstructured Data: An Overview. *Dataversity*, 2023. URL <https://www.dataversity.net/articles/structured-vs-unstructured-data-a-n-overview/>
- Christof Schoch. Big? Smart? Clean? Messy? Data in the Humanities. *Journal of Digital Humanities*, 2(3), 2023. URL <https://journalofdigitalhumanities.org/2-3/big-smart-clean-messy-data-in-the-humanities/>

September 2: Digital Platforms

Read:

- Clare Southerton. Datafication. In *The Encyclopedia of Big Data*, pages 358–361. Springer, 2022. URL https://doi.org/10.1007/978-3-319-32001-4_332-1
- José van Dijck, Thomas Poell, and and Martijn de Waal. *The Platform Society: Public Values in a Connective World*. Oxford University Press, Oxford, New York, November 2018. ISBN 978-0-19-088977-7, Chapter 2: “Platform Mechanisms.”

September 4: Digital Infrastructure

Read:

- Nicol Turner Lee and Darrell West. The Future of Data Centers, November 2025. URL <https://www.brookings.edu/articles/the-future-of-data-centers/>

Exercise: Representing Social Data.

Assignment: From Unstructured to Structured.

September 7: NO CLASS – Labor Day

September 9: Algorithms

Read:

- Taina Bucher. The algorithmic imaginary: Exploring the ordinary affects of Facebook algorithms. *Information, Communication & Society*, 20(1):30–44, January 2017. URL <https://doi.org/10.1080/1369118X.2016.1154086> (**read quickly**)
- Fanpage Karma Blog. The TikTok Algorithm Update: How it Works in 2025, 2025. URL <https://www.fanpagekarma.com/insights/the-2025-tiktok-algorithm-what-you-need-to-know/>

Exercise: From Platform to Pipeline.

Assignment: Platform Data Audit.

September 11: NO CLASS

September 14-18: Applied Lab One: Data Visualization

Read: No readings assigned.

Exercise: In-class lab (three days).

II. AI

September 21-25: Introduction to the AI Pipeline

Read:

- Melissa Heikkilä and Stephanie Arnett. This is where the data to build AI comes from. *MIT Technology Review.com*, December 2024. URL <https://www.technologyreview.com/2024/12/18/1108796/this-is-where-the-data-to-build-ai-comes-from/>
- Wendy M. Grossman. Ghost Work and Behind the Screen, book reviews: Lifting the veil on the internet's secret employment sector. *ZDNet Business*, 2019. URL <https://www.zdnet.com/article/ghost-work-and-behind-the-screen-book-reviews-lifting-the-veil-on-the-internets-secret-employment-sector/>
- Josh Dzieza. Inside the AI Factory. *Intelligencer*, June 2023. URL <https://nymag.com/intelligencer/article/ai-artificial-intelligence-humans-technology-business-factory.html>
- Kate Crawford and Vladan Joler. Anatomy of an AI System, 2018. URL <https://anatomyof.ai/>

Exercise: Build an AI Pipeline.

Assignment: Labeling as Interpretation.

September 28-October 2: How AI “Learns”

Read:

- Pedro Domingos. A Few Useful Things to Know About Machine Learning. *Communications of the ACM*, 55(10):78–87, October 2012. URL <https://dl.acm.org/doi/10.1145/2347736.2347755>

- Sebastian Raschka. Blog Post: Introduction to Machine Learning and Deep Learning, August 2020. URL <https://sebastianraschka.com/blog/2020/intro-to-dl-ch01.html>
- Jay Alammar. The Illustrated BERT, ELMo, and co. (How NLP Cracked Transfer Learning), 2021. URL <https://jalammar.github.io/illustrated-bert/> (read quickly / skim, and/or watch the video)
- Vipul Lohani. How Transformers Changed AI Forever: The Story Behind Modern Language Models, July 2025. URL <https://medium.com/@vipullohani8/how-transformers-changed-ai-forever-the-story-behind-modern-language-models-403cc3aea1a2>

Exercise: Thinking (/ Learning) Like a Machine.

Assignment: Learning Paradigms.

Bonus readings! Descriptive pieces titled “What is AI, Anyway?” from 1987, 2018, 2020, 2023, 2024, 2024 again, 2024 *again*, and 2025.

October 5-9: AI Systems in Practice: Recommenders and Rankings

Read:

- Ashok Chandrashekar, Fernando Amat, Justin Basilico, and Tony Jebara. Artwork Personalization at Netflix. *Medium*, December 2017. URL <https://netflixtechblog.com/artwork-personalization-c589f074ad76>
- Adam Pasick. The Magic That Makes Spotify’s Discover Weekly Playlists So Damn Good. *Quartz*, July 2022. URL <https://qz.com/571007/the-magic-that-makes-spotifys-discover-weekly-playlists-so-damn-good>
- Ben Smith. How TikTok Reads Your Mind. *The New York Times*, December 2021. URL <https://www.nytimes.com/2021/12/05/business/media/tiktok-algorithm.html>
- Prem Vishnoi. Content-Based vs Collaborative Filtering: How TikTok and Netflix Hack Your Attention, August 2025. URL <https://medium.com/nextgenllm/content-based-vs-collaborative-filtering-how-tiktok-and-netflix-keep-you-addicted-75b444a09c01>

Exercise: Design a Recommender.

Assignment: Ranking Redesign.

October 12-14: Automated Conversation

Read:

- Niklas Heidlöff. Foundation Models, Transformers, BERT and GPT, February 2023. URL <https://heidloff.net/article/foundation-models-transformers-bert-and-gpt/>

- Murray Shanahan. Talking about Large Language Models. *Communications of the ACM*, 67(2):68–79, February 2024. URL <https://cacm.acm.org/research/talking-about-large-language-models/>
- Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. On the Dangers of Stochastic Parrots: Can Language Models Be Too Big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, FAccT '21, pages 610–623, New York, NY, USA, March 2021. Association for Computing Machinery. ISBN 978-1-4503-8309-7. URL <https://dl.acm.org/doi/10.1145/3442188.3445922>
- Jamie Bartlett. Meet the AI Jailbreakers: ‘I See the Worst Things Humanity Has Produced’. *The Guardian*, April 2026. ISSN 0261-3077. URL <https://www.theguardian.com/technology/2026/apr/29/meet-the-ai-jailbreakers-i-see-the-worst-things-humanity-has-produced>
- Optional: Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention Is All You Need. In *31st Conference on Neural Information Processing Systems (NIPS 2017)*. arXiv, August 2023. URL <http://arxiv.org/abs/1706.03762> (the original “Attention is all you need” paper)
- Also optional: Stephen Wolfram. What Is ChatGPT Doing ... and Why Does It Work?, February 2023. URL <https://writings.stephenwolfram.com/2023/02/what-is-chatgpt-doing-and-why-does-it-work/> (link only – a detailed blog post by polymath Stephen Wolfram on how neural nets and LLMs work, for those who want to nerd out)

October 16: NO CLASS

Exercise: None

Assignment: Stress-Testing The Parrot.

October 19-21: Generative AI

Read:

- Chris Tozzi. Generative Models: VAEs, GANs, Diffusion, Transformers, NeRFs, 2025. URL <https://www.computerweekly.com/de/tipp/Generative-Modelle-VAE-GAN-Diffusion-Transformer-NeRF>
- Amazon Web Services. What is RAG? - Retrieval-Augmented Generation AI Explained - AWS, 2026. URL <https://aws.amazon.com/what-is/retrieval-augmented-generation/>

(Read these three quickly:)

- Bella Isaacs-Thomas. How AI Turns Text Into Images. *PBS News*, January 2023. URL <https://www.pbs.org/newshour/science/how-ai-makes-images-based-on-a-few-words>
- Jewly Hight. AI Music Isn’t Going Away. Here Are Four Big Questions About What’s Next. *NPR*, April 2024. URL <https://www.npr.org/2024/04/25/1246928162/generative-ai-music-law-technology>

- Faruk Alpay. Same AI, Different Answer: How Tiny Prompts Can Change Everything, August 2025. URL <https://lightcapai.medium.com/same-ai-different-answer-how-tiny-prompts-can-change-everything-83e880f9773f>

Exercise: Real Or Not?

October 23: Agentic AI

Read:

- Maria Mouschoutzi. AI Agents Explained: What Is a ReAct Loop and How Does It Work?, July 2026a. URL <https://towardsdatascience.com/ai-agents-explained-what-is-a-react-loop-and-how-does-it-work/> (Consider also reading Mouschoutzi's preceding / lead-in post: Maria Mouschoutzi. Tool Calling, Explained: How AI Agents Decide What to Do Next, June 2026b. URL <https://towardsdatascience.com/tool-calling-explained-how-ai-agents-decide-what-to-do-next/>)
- Ethan Mollick. Real AI Agents and Real Work, September 2025. URL <https://www.oneusefulthing.org/p/real-ai-agents-and-real-work>

Assignment: Who Made This?

October 26-30: Applied Lab Two: Tuning AI Systems With Data

Read: No readings assigned.

Exercise: In-class lab (three days).

III. Society

November 2-6: Surveillance, Privacy, and Digital Consent

Read:

- Jim Isaak and Mina J. Hanna. User Data Privacy: Facebook, Cambridge Analytica, and Privacy Protection. *Computer*, 51(8):56–59, August 2018. URL <https://ieeexplore.ieee.org/document/8436400/>
- On regulation:
 - Arthur Beesley. Meta fined record €1.2bn by Irish regulator for violating European privacy rules. *The Irish Times*, May 2023. URL <https://www.irishtimes.com/technology/big-tech/2023/05/22/facebook-owner-meta-fined-record-12bn-by-irish-regulator-for-violating-european-privacy-rules/>

- Kilian Vieth-Ditlmann. A Guide to the AI Act, the EU's New AI Rulebook, June 2025. URL <https://algorithmwatch.org/en/ai-act-explained/>
- Miriam Lord. US Copyright Office on AI: Human Creativity Still Matters, Legally, 2025. URL <https://www.wipo.int/en/web/wipo-magazine/article-details/?assetRef=73696&title=us-copyright-office-on-ai-human-creativity-still-matters-legally>
- Shoshana Zuboff. Chapter One. In *The Age of Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power*. Public Affairs, New York, NY, first edition edition, 2019
- Nikki Goth Itoi. Be Careful What You Tell Your AI Chatbot. *Stanford HAI*, 2025. URL <https://hai.stanford.edu/news/be-careful-what-you-tell-your-ai-chatbot>

Exercise: Read the Fine Print.

Assignment: A Tale of Two Frameworks.

November 9-13: Bias, Inequality, and Algorithmic Harm

Read:

- Arvind Narayanan and Sayash Kapoor. How Predictive AI Goes Wrong. In *AI Snake Oil: What Artificial Intelligence Can Do, What It Can't, and How to Tell the Difference*, pages 36–59. Princeton University Press, 2024
- Julia Angwin, Jeff Larson, Lauren Kirchner, and Surya Mattu. Machine Bias, May 2016. URL <https://www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing>
- Carolyn Johnson. Racial Bias in a Medical Algorithm Favors White Patients Over Sicker Black Patients. *The Washington Post*, October 2019. URL <https://www.washingtonpost.com/health/2019/10/24/racial-bias-medical-algorithm-favors-white-patients-over-sicker-black-patients/>
- Stefan Milne. AI Tools Show Biases in Ranking Job Applicants' Names According to Perceived Race and Gender. *UW News*, October 2024. URL <https://www.washington.edu/news/2024/10/31/ai-bias-resume-screening-race-gender/>
- Sujata Gupta. Biased Online Images Train AI Bots to See Women as Younger. *ScienceNews*, October 2025. URL <https://www.sciencenews.org/article/bias-online-images-women-ai-hiring>

Exercise: Pick Your Fairness.

Assignment: Diagnose and Prescribe.

November 16-20: Environmental Impacts

Read:

- Michelle Kim. 10 Breakthrough Technologies 2026: Hyperscale AI data centers. *MIT Technology Review*, 2026. URL <https://www.technologyreview.com/2026/01/12/1129982/hyperscale-ai-data-centers-energy-usage-2026-breakthrough-technology/>

- Eli Tan and Justin Chambers. Their Water Taps Ran Dry When Meta Built Next Door. *The New York Times*, July 2025. ISSN 0362-4331. URL <https://www.nytimes.com/2025/07/14/technology/meta-data-center-water.html>
- Katherine Bourzac. Generative AI Has a Massive E-Waste Problem. *IEEE Spectrum*, November 2024. URL <https://spectrum.ieee.org/e-waste>
- Henry Epp. Data centers lowered electric bills in some places — for now. *Marketplace*, July 2026. URL <https://www.marketplace.org/story/2026/07/10/data-centers-lowered-electric-bills-in-some-places-for-now>
- Rachel Mural. How local communities are challenging Big Tech data centers’ noise, pollution and rising electricity bills. *The Conversation*, June 2026. doi: 10.64628/AAI.jme9vxdqx. URL <https://theconversation.com/how-local-communities-are-challenging-big-tech-data-centers-noise-pollution-and-rising-electricity-bills-284642>
- Optional / Recommended: James O’Donnell and Casey Crownhart. We did the math on AI’s energy footprint. Here’s the story you haven’t heard. *MIT Technology Review*, May 2025. URL <https://www.technologyreview.com/2025/05/20/1116327/ai-energy-usage-climate-footprint-big-tech/> (access link [here](#))

Exercise: What Did That Cost?

Assignment: The Ground Truth.

November 23-27: No Class – Thanksgiving Break

November 30 - December 3: AI’s Impact on Society, Culture, and Labor

Read:

- Ziv Epstein, Aaron Hertzmann, and The Investigators of Human Creativity. Art and the Science of Generative AI. *Science*, 380(6650):1110–1111, June 2023. URL <https://www.science.org/doi/10.1126/science.adh4451>
- Cory Doctorow. AI Companies Will Fail. We Can Salvage Something From The Wreckage. *The Guardian*, January 2026. URL <https://www.theguardian.com/us-news/ng-interactive/2026/jan/18/tech-ai-bubble-burst-reverse-centaur>
- Matt Shumer. Something Big Is Happening, February 2026. URL <https://shumer.dev/something-big-is-happening>
- Samantha Oltman. The Rise of AI is Making the Future of Work Look Bleak – But It Could Be an Opportunity. *The Guardian*, February 2026. URL <https://www.theguardian.com/technology/ng-interactive/2026/feb/19/ai-work-future>
- Kyle Saunders. AIPolitics 27: Republicans Told AI Companies Data Centers Are Poison. Shapiro Signed Their Recommended Policy the Same Day, August 2026. URL <https://kylesaunders.substack.com/p/aipolitics-27-republicans-told-ai>

- Ion Valis. Issue 84 | The Attention Economy is Dead. Welcome to the Attention Society., March 2026. URL <https://ionintelligence.substack.com/p/issue-84-the-attention-economy-is>

Exercise: Centaur or Reverse Centaur?

Assignment: Interview with an AI.

December 7-11: Governance and Responsible AI

Read:

- Michelle Faverio and Emma Kikuchi. What the Data Says About Americans' Views of Artificial Intelligence, March 2026. URL <https://www.pewresearch.org/?p=7828>
- Greg Rosalsky. Why the AI World is Suddenly Obsessed with a 160-Year-Old Economics Paradox. *NPR's Planet Money*, February 2025. URL <https://www.npr.org/sections/planet-money/2025/02/04/g-s1-46018/ai-deepseek-economics-jevons-paradox>
- Nicole Greenfield. AI Data Centers: Big Tech's Impact on Electric Bills, Water, and More. *Consumer Reports*, March 2026. URL <https://www.consumerreports.org/data-centers/ai-data-centers-impact-on-electric-bills-water-and-more-a1040338678/> (Note: This is a much nicer read at the URL than in its PDF form.)
- Nick Robins-Early. AI Companies Know They Have an Image Problem. Will Funding Policy Papers and Thinktanks Dig Them Out? *The Guardian*, April 2026. ISSN 0261-3077. URL <https://www.theguardian.com/technology/2026/apr/12/ai-image-problem-policy-papers-thinktanks>
- Optional: Dario Amodei. The Adolescence of Technology, January 2026. URL <https://www.darioamodei.com/essay/the-adolescence-of-technology> (Long, but readable; focus on "Player Piano" and "Humanity's test" if you don't want to read the whole thing.)

Exercise: The AI Constitution.

Assignment: None.

December TBA, 2026: Final Examination