

Prompt Engineering +

Large Language Models +

ChatGPT +

Ai Bias +

Artificial Intelligence +


Same AI, Different Answer: How Tiny Prompts Can Change Everything



Faruk Alpay

Follow

12 min read · Aug 13, 2025



Why Does ChatGPT Sometimes Feel Different?

If you've used AI chatbots like ChatGPT for a while, you may have noticed something odd: ask the **same** question twice, or on different days or devices, and the answers can vary. One day it cracks a joke, the next day it's all business. You're not imagining things — AI responses really can shift like quicksand under your feet. But why does this happen? The short answer is

that **context is king**. These AI models are extremely sensitive to how you ask, hidden system settings, and even periodic updates. A slight tweak in wording or environment can lead to a surprisingly different outcome — a bit like a digital butterfly effect where a tiny change in the prompt can stir up a storm in the response.

Part of the reason is that modern AI assistants are **trained to be helpful and safe**, which sometimes makes them behave cautiously or neutrally by default. OpenAI's **ChatGPT** (based on GPT-4) was fine-tuned with human feedback to follow instructions and avoid toxic or untruthful outputs. That alignment process greatly improved the quality and safety of answers, but also meant the AI learned certain preferences — for example, to refrain from edgy humor or controversial takes unless explicitly asked. So if it feels like ChatGPT has “lost its spark” or is too neutral, that's often by design. It's a bit like an overly polite coworker: it defaults to careful professionalism, unless you encourage a more casual tone.

The Butterfly Effect of Prompt Wording

It turns out that **tiny changes in how you phrase your prompt can have outsized effects on an AI's answer**. Researchers have aptly called this the “*butterfly effect*” of prompting. For instance, simply adding an extra space at the end of your question, or saying “Thanks” at the conclusion, might flip the result you get. It sounds absurd, but under the hood these models are probabilistic word jugglers. They predict the next word based on probabilities, and a minor prompt perturbation can nudge those probabilities in a different direction. A recent study demonstrated this vividly: **even the smallest prompt perturbation (like an innocuous whitespace or punctuation change) sometimes made an LLM change its answer entirely**. In other cases, asking for a specific *format* — say, “Give the answer in JSON” — can push the model into a different mode, influencing what content it emphasizes or skips.

Why are they so sensitive? Think of the AI's knowledge as a huge map of patterns. Your exact phrasing determines which path the model takes through that map. If you start your question with “*Explain why...*” versus “*Give a brief explanation why...*”, the model might traverse a slightly different route: one path might lead it to a concise textbook-style answer, while

another invites a longer, more narrative explanation. This isn't just theory; tests with GPT-4 have shown that **different prompts lead to responses of varying quality and length**, even for the same underlying question. In a healthcare Q&A study, for example, adding more detailed instructions in the prompt (telling the AI its role or the desired depth) produced much longer answers — though not always *better* answers. The key lesson is that the AI is highly literal: it tries to do exactly what you ask, and even subtle wording cues signal it to respond in a certain style or scope.

This sensitivity also explains those notorious “**jailbreak**” **tricks** you may have heard of. By phrasing a request in a clever way, users can sometimes get an AI to bypass its usual filters. The research community has found that certain prompt variations or exploits can **dramatically alter an AI's behavior**, essentially tricking it to ignore rules. For example, inserting a phrase like “*begin roleplay mode*” or an HTML tag request might sneak past content safeguards and yield an answer the AI normally wouldn't give. From the AI's perspective, these are just alternate routes in that map of patterns — ones that happen to lead off the safe, well-lit paths. It's a cat-and-mouse game: as companies patch the obvious tricks, new prompt tweaks emerge. The takeaway for regular users is that phrasing matters not just for *what* answer you get, but sometimes for *whether* you get an answer at all.

Under the Hood: Context and “Hidden” Instructions

Another reason the same AI can behave differently is that there's more in the conversation than just your typed question. These AI systems operate with **layers of context**. OpenAI's ChatGPT, for instance, doesn't start from a blank slate each time — it has system-level instructions (often invisible to the user) and memory of the dialogue. If those differ, the responses differ. An interesting case study came from ChatGPT's own official apps: users noticed that **ChatGPT on mobile tended to give shorter, snappier answers**, whereas on the web it was more verbose. This wasn't the AI “having a mood swing” so much as a design choice. The mobile apps included a hidden instruction to keep answers brief (better for on-the-go reading), while the web interface allowed more elaborate replies. In short, the **system prompt** — an initial guideline given to the model — was different on iPhone vs. desktop. The result? Ask “*Why is the sky blue?*” on your phone and you'd get a two-sentence answer; ask the exact same thing on your laptop and you might get a detailed

paragraph with a mini science lecture. The AI was the same, but the ground rules it was given differed.

Likewise, conversation history plays a big role. These models have a limited “memory” (context window) of the recent dialogue. If you’ve been chatting at length, earlier parts of the conversation can bias later answers. A question asked cold in a new chat might get a different answer than the same question asked after a long discussion, because the AI is taking into account everything said so far. For instance, if earlier you told the AI that you prefer informal language, it will likely remember that and respond more casually later. This context sensitivity means that to reproduce an answer consistently, one often must reproduce the surrounding context *exactly*. Even something as small as a prior user message saying “I’m just brainstorming here...” could nudge the tone of the AI’s next answer (perhaps making it more exploratory or cautious).

In essence, **there is always an invisible conversation happening alongside the visible one**. The AI’s “decision-making” is heavily influenced by these hidden or earlier inputs. Companies like Anthropic have even made those guiding principles explicit in training their models. Anthropic’s Claude AI is built on a concept called “**Constitutional AI**,” where the system follows a set of written principles (a kind of AI constitution) to decide what is appropriate. The result is a model that tries to be *harmless yet helpful* by internally critiquing its outputs against those rules. Notably, this leads to specific behavioral differences — for example, Claude will engage with a sensitive or potentially harmful query but then explain its objections or concerns, rather than just refusing with a generic apology. The **principles in its prompt** guide it to handle tough questions by *explaining* why it won’t do something, whereas another AI without those might simply say “No, I can’t.” So if you ask an ethically tricky question to Claude and to ChatGPT, you may notice Claude gives a more nuanced mini-lecture on safety, while ChatGPT might respond with a shorter refusal. They’re both aligned AIs, but their “rulebooks” (and thus their hidden system prompts) aren’t identical.

Model Upgrades: When AI “Changes” Overnight

Apart from prompt phrasing and context, there’s another reason your AI’s tone or quality might seem different over time: **the AI itself is a moving**

target. These systems get updated and improved, often behind the scenes. OpenAI periodically pushes new versions of ChatGPT (or entirely new models like GPT-4), and those updates can **change the behavior** in noticeable ways. A striking real-world example occurred with GPT-4 from March to June of 2023. Users were surprised to find that GPT-4 in summer 2023 felt less capable in some areas than it had been initially. Indeed, a study confirmed that **the very same GPT-4 service had significant performance drifts over a few months.** In March, it was acing questions about prime numbers with ~84% accuracy, but by June, on the same questions, accuracy dropped to ~51%. It also became more hesitant to answer certain sensitive questions and was more likely to follow the letter of user instructions rather than the spirit. OpenAI hadn't announced a downgrade — rather they likely fine-tuned the model to be safer or fixed bugs, which inadvertently altered some behaviors. The net effect for users was that **ChatGPT “felt” different** — enough that some started complaining it had been “nerfed” or dumbed down.

This phenomenon isn't unique to OpenAI. AI models are improving rapidly, and with improvement comes change. When **Anthropic upgraded Claude to Claude 2**, they expanded its memory and made it better at detailed reasoning. Many users found Claude 2's answers longer and more comprehensive by default, because the new model had been trained to utilize that broader context (up to 100K tokens) and go deeper into analysis. Similarly, Google's **Gemini**, a next-generation model announced by DeepMind, is said to have capabilities that *far* surpass earlier models on many benchmarks. Gemini was built from the ground up to incorporate advanced reasoning — it can “think through” problems internally before responding. If you switch from a model that answers in one go to a model like Gemini that can deliberate, you'll notice differences. For example, Gemini might take an extra moment but give a more structured and accurate answer, whereas a simpler model might blurt the first thing that comes to mind (or even make a mistake it could have caught by thinking longer). In Google's tests, the largest version, **Gemini Ultra**, even exceeded human-level performance on tough academic benchmarks by carefully reasoning through questions. What this means for you: as these AI services update to new brains, you might suddenly observe your trusty chatbot responding with a new style. It could become more precise, more verbose, or in some cases *more reserved*, depending on the alignment and objectives of the update.

Crucially, companies often don't fully disclose every little tuning, so from the outside it can feel like the AI woke up on the wrong side of the bed one day. In reality, it's usually a subtle shift in the model's training or system prompts. The **behavior of the "same" AI can change substantially in a short time**, so if you're using these tools long-term, you almost have to *relearn* some quirks with each major version. The silver lining is that most changes aim to make the AI better — more factual, more capable, more aligned with user needs. The downside is occasional growing pains: you might have to adjust your prompts or approach when a new version doesn't respond like the old one did.

Steering the AI (Or: How to Get the Answer You Want)

The good news in all of this is that you, as the user, have a lot of influence over how an AI responds. Think of it less like a magic 8-ball that gives random answers, and more like a talented but slightly foreign assistant — it *wants* to give you what you're after, but you may need to show it the way. Here are a few tips to **mitigate prompt sensitivity and coax the best output**:

- **Be Explicit with Style or Tone:** If you want a casual, humorous answer, say so! Don't just ask "*Explain quantum physics.*" Try "*Explain quantum physics in a casual, humorous tone.*" These models will usually oblige. They have been trained on countless styles, from jokey to formal, and will mirror the tone you request or exemplify. In fact, the more you write in a certain style, the more the AI will subtly adapt to match it. So set the vibe early. As one AI observer quipped, using a new model is like **getting a new coworker** — the coworker is brilliant, but you need to **onboard** them to your preferred communication style. Tell the AI the role or persona you want it to take ("You are a friendly tutor...") and it will respond in kind.
- **Provide Examples or Formats:** If the first answer you get isn't what you hoped, consider guiding the model by example. Want a bulleted list instead of a paragraph? Follow your question with a short dummy example: "*Give me 3 tips (e.g., 'Tip 1: ...')*". The AI will pick up on the hint and format accordingly. Similarly, if the answer was too shallow, you can push it towards depth by asking a more specific follow-up or by framing the question with context ("*Explain like I'm a programmer curious about the math behind it.*"). Because the AI's output is so prompt-dependent, a well-

crafted prompt is often the difference between a generic answer and an excellent one. Entire guides on “prompt engineering” have emerged for this reason.

- **Use Custom Instructions or Memory Wisely:** Some platforms (like ChatGPT) let you set persistent preferences (e.g. “Always respond with metric units” or “Keep answers under 3 paragraphs unless I say otherwise”). Take advantage of these features if available — they act as a personalized system prompt for you. Over a long session, if the AI starts veering off course or getting too terse, gently remind it of the context: “*You can elaborate more,*” or “*Remember, we’re discussing a lighthearted take here.*” The models do try to follow your lead, and consistency helps. If you maintain a certain framing across questions, the AI is more likely to stay in the zone. On the flip side, if you notice it going in an unwanted direction (say, getting too technical), you can course-correct with a follow-up instruction.
- **Mind the Temperature (if available):** Some AI interfaces allow you to adjust a “*randomness*” setting (often called *temperature*). A higher temperature means the AI is more likely to take risks or generate varied answers — great for creativity, but it can also cause inconsistency. A lower temperature makes it more deterministic and focused (it will stick to its highest-probability answer each time, reducing variability). If you felt the AI gave you one good answer and one weird one for the same prompt, a high temperature might be why. Tuning this setting down can make responses more stable. Conversely, if you want it to brainstorm lots of different ideas, turning it up will do that (but expect more variability in quality).

In short, **you have levers to pull.** A large language model can be seen as an instrument, and your prompt is how you play it. By understanding its sensitivities, you can get it to sing in the key you desire.

The Road Ahead: More Stable, More Personal AIs

It’s important to remember that we’re still in the early days of mainstream AI assistants. Many of these quirks — prompt sensitivity, variability, unexpected changes — are active areas of research and engineering. The industry is well aware that unpredictability is a problem. In fact, one reason **Google**


DeepMind developed Gemini with an emphasis on reasoning is to tackle exactly this: make the model think more systematically, so that it's less likely to veer off due to a random prompt tweak. By having the AI internally plan out its answers (a bit more like how a human expert might), the hope is that it won't be so easily thrown off by phrasing. Other research efforts focus on **prompt robustness** — training models on many rephrasings of the same queries so they learn to give consistent outputs. There's even the idea of *prompt ensembling*, where multiple prompts are tried and the answers combined, to cancel out individual prompt bias. All of this is aimed at a future where you don't have to worry about crafting the *perfect* question to get a good answer — the AI will handle imperfect prompts gracefully.

At the same time, AI developers are giving users more ways to *shape* the AI's behavior deliberately. OpenAI, for example, introduced custom instructions for ChatGPT, allowing you to set your own defaults (like always answer as a Socratic mentor, or always provide an analogy). This is a recognition that one size doesn't fit all — some folks want a playful assistant, others want a terse encyclopedia. In the future, we might see AI that automatically personalizes itself to your style over time, effectively learning your preferences. Imagine never having to say “in a funny tone” because the AI already knows you love humor. We're heading in that direction.

However, even as consistency improves, **change will be a constant**. AI models will continue to evolve rapidly. A year from now, the hot new model (perhaps GPT-5 or Gemini Ultra 2) might make today's best look quaint. With each leap, there will be new capabilities and new tuning that we'll adapt to. It's a bit like software updates — mostly beneficial, but occasionally moving your cheese. The difference is your “software” here is an intelligent entity that converses with you, so you feel the changes very viscerally. The best mindset is to stay curious and flexible. When your AI feels different, ask “What might have changed? Did I ask differently, or is it a new model behavior?” — because as we've seen, both can happen.

Bottom Line: You're not crazy if ChatGPT or another AI suddenly behaves differently. There are concrete reasons behind the magic, from prompt wording to hidden instructions to model upgrades. The exciting part is that once you *know* this, using AI becomes less of a shot in the dark. You gain the ability to debug and direct the conversation, making these tools work better

for you. In a way, understanding an AI's fickleness is actually empowering — it turns mystery into mastery. And as these systems continue to improve, we can look forward to a future where talking to an AI feels more and more like a natural, reliable dialog — **with fewer unwelcome surprises, and yet still full of delightful creativity.**



One comma, two different answers.

Prompt Engineering +

Large Language Models +

ChatGPT +

Ai Bias +

Artificial Intelligence +

Some rights reserved ⓘ



Written by Faruk Alpay

190 followers · 12 following

Follow

Faruk Alpay is a Data Engineer and founder of Lightcap, known for discovering Alpay Algebra and exploring the deep links between math, AI, and identity.

No responses yet

